

# Scalability and Distributed Systems

Bill Howe

INFX 575: Data Science III

Scaling, Applications, and Ethics

## This week...

- Quiz
- Scalability
- MapReduce
- Class Exercise: Launching a Spark cluster
- Twitter Assignment due tonight
- Project Proposal due next week


<https://pollev.com/billhowe>


# What Does Scalable Mean?

- Operationally:
  - In the past: “Works even if data doesn’t fit in main memory”
  - Now: “Can make use of 1000s of cheap computers”
- Algorithmically:
  - In the past: If you have  $N$  data items, you must do no more than  $N^m$  operations -- “polynomial time algorithms”
  - Now: If you have  $N$  data items, you must do no more than  $N^m/k$  operations, for some large  $k$ 
    - Polynomial-time algorithms must be parallelized
  - Soon: If you have  $N$  data items, you should do no more than  $N * \log(N)$  operations
    - As data sizes go up, you may only get one pass at the data
    - The data is streaming -- you better make that one pass count
    - Ex: Large Synoptic Survey Telescope (30TB / night)

## Example: Find matching DNA sequences

- Given a set of sequences
- Find all sequences equal to “GATTACGATATTA”






TACCTGCCGTAA = GATTACGATATTA?

No.

time = 0




CCCCCAATGAC = GATTACGATATTA?

No.

time = 1






GATTACGATATTA contains GATTACGATATTA?

Yes!

Send it to the output.


time = 17



40 records, 40 comparisons

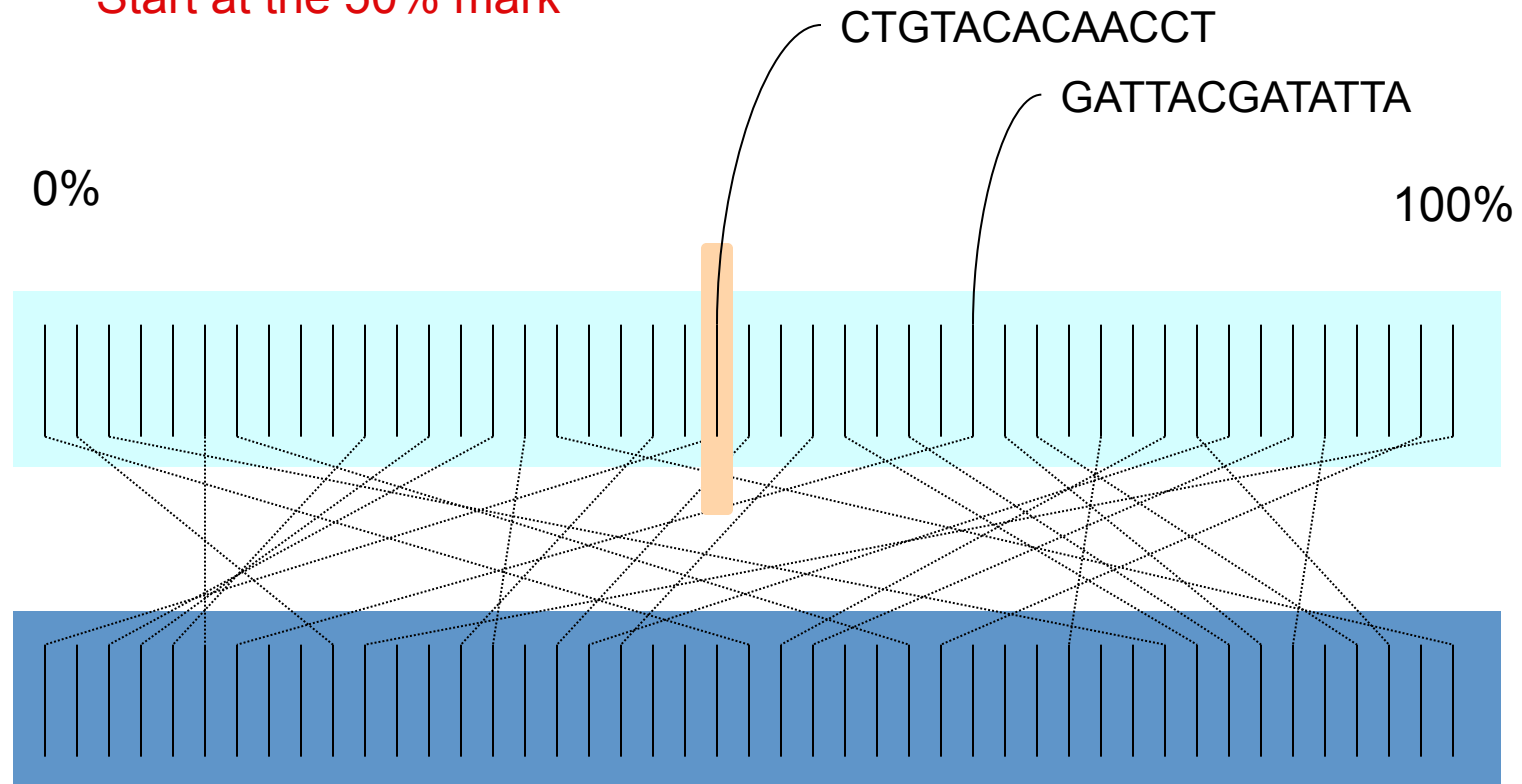
$N$  records,  $N$  comparisons

The algorithmic complexity is order  $N$ :  $O(N)$



What if we sort the sequences?

Start at the 50% mark




CTGTACACAACCT < GATTACGATATTA

time = 0

No match.


Skip to 75% mark



time = 1


No match.

Go back to 62.5% mark



No match.


Skip back to 68.75% mark



GATTACGATATTA = GATTACGATATTA

Match!

Walk through the records until we fail to match.



How many comparisons did we do?

40 records, only 4 comparisons

$N$  records,  $\log(N)$  comparisons

This algorithm is  $O(\log(N))$  Far better scalability



# Relational Databases

- Databases are good at “Needle in Haystack” problems:
  - Extracting small results from big datasets
  - Transparently provide “old style” scalability
  - Your query will **always**\* finish, regardless of dataset size.
- Indexes are easily built and automatically used when appropriate


```
CREATE INDEX seq_idx ON sequence(seq);
```


```
SELECT seq  
  FROM sequence  
 WHERE seq = 'GATTACGATATTA';
```

**\*almost**

## New task: Read Trimming


- Given a set of DNA sequences
- Trim the final  $n$  bps of each sequence
- Generate a new dataset






TACCTGCCGTAA becomes TACCT

time = 0




CCCCCAATGAC becomes CCCCC

time = 1



GATTACGATATTA becomes GATTA

time = 17





Can we use an index?


No. We have to touch every record no matter what.

The task is fundamentally  $O(N)$


Can we do any better?









time = 0




time = 1



time = 2



time = 3



How much time did this take?

7 cycles

40 records, 6 workers

time = 7

$O(N/k)$

Poll here:

<https://pollev.com/billhowe>